

hw2

Apoorva Srinivasan

10/16/2019

Question 4

```
SAheart = read.table("http://www-stat.stanford.edu/~tibs/ElemStatLearn/datasets/SAheart.data", sep=",", header=TRUE)
```

```
train = SAheart[1:300,]  
test = SAheart[301:462,]
```

```
log_model = glm(chd ~ ., family = binomial(link = "logit"), data = train)  
summary(log_model)
```

```
##  
## Call:  
## glm(formula = chd ~ ., family = binomial(link = "logit"), data = train)  
##  
## Deviance Residuals:  
##      Min       1Q   Median       3Q      Max   
## -1.8176  -0.8282  -0.4348   0.9484   2.4491   
##  
## Coefficients:  
##              Estimate Std. Error z value Pr(>|z|)      
## (Intercept)  -5.013991   1.570094  -3.193   0.00141 **   
## sbp          -0.004784   0.007500  -0.638   0.52356      
## tobacco       0.069165   0.032321   2.140   0.03236 *     
## ldl           0.111462   0.074524   1.496   0.13474      
## adiposity     0.046678   0.037295   1.252   0.21072      
## famhistPresent 0.826525   0.280107   2.951   0.00317 **   
## typea         0.046926   0.015911   2.949   0.00319 **   
## obesity      -0.063287   0.053181  -1.190   0.23404      
## alcohol       0.005111   0.005927   0.862   0.38849      
## age           0.039145   0.014631   2.675   0.00746 **   
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## (Dispersion parameter for binomial family taken to be 1)  
##  
##      Null deviance: 394.29  on 299  degrees of freedom  
## Residual deviance: 317.75  on 290  degrees of freedom  
## AIC: 337.75  
##  
## Number of Fisher Scoring iterations: 4  
  
fitted.results = predict(log_model, test, type = "response")  
fitted.results <- ifelse(fitted.results > 0.5, 1, 0) #picking 0.5 as the boundry  
log_error = mean((test$chd - fitted.results)^2)  
log_error  
  
## [1] 0.2530864
```

```
##same as log_error = mean(fitted.results != test$chd)

std_log_error = sd((test$chd - fitted.results)^2)/ sqrt(nrow(test))
std_log_error

## [1] 0.03426547
```

LDA

```
lda_model = lda(chd~., data = train)
summary(lda_model)
```

```
##          Length Class  Mode
## prior      2      -none- numeric
## counts     2      -none- numeric
## means     18      -none- numeric
## scaling    9      -none- numeric
## lev        2      -none- character
## svd         1      -none- numeric
## N           1      -none- numeric
## call        3      -none- call
## terms       3      terms  call
## xlevels     1      -none- list
```

```
lda_pred = predict(lda_model, test)$class
lda_error = mean((test$chd - lda_pred)^2)
```

```
## Warning in Ops.factor(test$chd, lda_pred): '-' not meaningful for factors
lda_error
```

```
## [1] NA

lda_pred = as.numeric(lda_pred)
test$chd = as.numeric(test$chd)
std_lda_error = sd((test$chd - lda_pred)^2)/ sqrt(nrow(test))
std_lda_error

## [1] 0.08431974
```

QDA

```
qda_model = qda(chd~., data = train)
qda_model$means
```

```
##          sbp  tobacco      ldl adiposity famhistPresent      typea  obesity
## 0 134.0105 2.629368 4.468421 23.29147      0.2947368 53.04211 25.57868
## 1 138.9091 5.433818 5.451545 27.90682      0.5545455 56.28182 26.76291
##      alcohol      age
## 0 12.41268 38.02632
## 1 18.75418 49.47273
```

```
qda_pred = predict(qda_model, test)$class
qda_pred = as.numeric(qda_pred)
```

```

qda_error = mean((qda_pred - test$chd)^2)
std_qda_error = sd((qda_pred - test$chd)^2)/sqrt(nrow(test))

log = cbind("Logistic Regression", log_error, std_log_error)
lda = cbind("LDA", lda_error, std_lda_error)
qda = cbind("QDA", qda_error, std_qda_error)

summary_table = rbind(log, lda, qda)
colnames(summary_table) = c("Model", "Test Error", "Std Error")

```

Given that the test error and the standard errors are similar for all three models, I'd pick logistic since it's the easiest to interpret.

Question 3